

Resampling

Contents

4.3 Resampling	2
Recall - Resampling from the Population	2
Mimic Resampling	3
The Bootstrap Method	4
Shark Lengths Example	5
An estimate of the sampling distribution	6
Shark Length Comparision	6
Bootstrap Standard Deviation	8
The Average	9
Comparing two estimators of the standard deviation	9
Aside: n Versus $n - 1$	11
Estimating the sampling bias	12
Biased Corrected Estimators	12
Example with different Attributes	13
4.5.2 Bootstrap confidence intervals	14
Estimating the standard bootstrap interval coverage probability	15
4.5.2.2 Percentile method	16
Bootstrap versus Sampling Distribution	16
Percentile Method	17
Assessing the coverage probability	17
4.5.2.1 Bootstrap-t confidence intervals	19
Sampling Distribution Approximation	19
The bootstrap- t	20
Sampling Distribution Comparison	21
Bootstrap- t Coverage probability	22
A thought on Standard Deviations	22
A Double Double	23
Example	24
Assessing the coverage probability	25
Comments	25

- We will use the following throughout the slide-set.

```
sharks <- read.csv("../Data/Sharks/sharks.csv")
popSharks <- rownames(sharks)

# A function to put n or N in the denominator of SD
# rather than n-1 or N-1
sdn <- function( y.pop ) {
  N = length(y.pop)
  sqrt( var(y.pop)*(N-1)/(N) ) }

#samples <- combn(popSharks, 5)
#N_s <- ncol(samples)
```

```

N_s <- 10^4
n = 6
set.seed(341)
samples <- sapply(1:N_s, FUN =function(b) sample(popSharks, n, replace = TRUE) )

avePop <- mean(sharks[, "Length"])
avesSamp <- apply(samples, MARGIN = 2,
                  FUN = function(s){mean(sharks[s, "Length"])})
sampleErrors <- avesSamp - avePop

tmpAve <- mean(avesSamp)
tmpSD <- sdn(avesSamp)

sdsSamp <- apply(samples, MARGIN = 2,
                  FUN = function(s){sdn(sharks[s, "Length"])})

```

4.3 Resampling

- As shown in previous sections, understanding the sampling behaviour of sample attributes is essential for making inferences about any population attribute. e.g.
 - For a discrepancy measure whose sampling distribution allows us to test hypotheses and
 - a pivotal quantity whose sampling distribution allows us to construct confidence intervals.
- However, this process requires undertaking repeated sampling from the population
- In practice, however,
 - the population from which our sample was taken cannot be repeatedly sampled for our purposes,
 - we have only one sample.

Recall - Resampling from the Population

- Draw sample \mathcal{S} of size n from a study population \mathcal{P} according to some sampling mechanism
 - then calculate the sample attribute $a(\mathcal{S})$ to estimate its population counterpart $a(\mathcal{P})$.
- To understand the sampling distribution of any attribute $a(\mathcal{S})$,
 - we draw M samples $\mathcal{S}_1, \dots, \mathcal{S}_M$ and

- use the values $a(\mathcal{S}_1), \dots, a(\mathcal{S}_M)$ to inform us about the sampling distribution of $a(\mathcal{S})$.

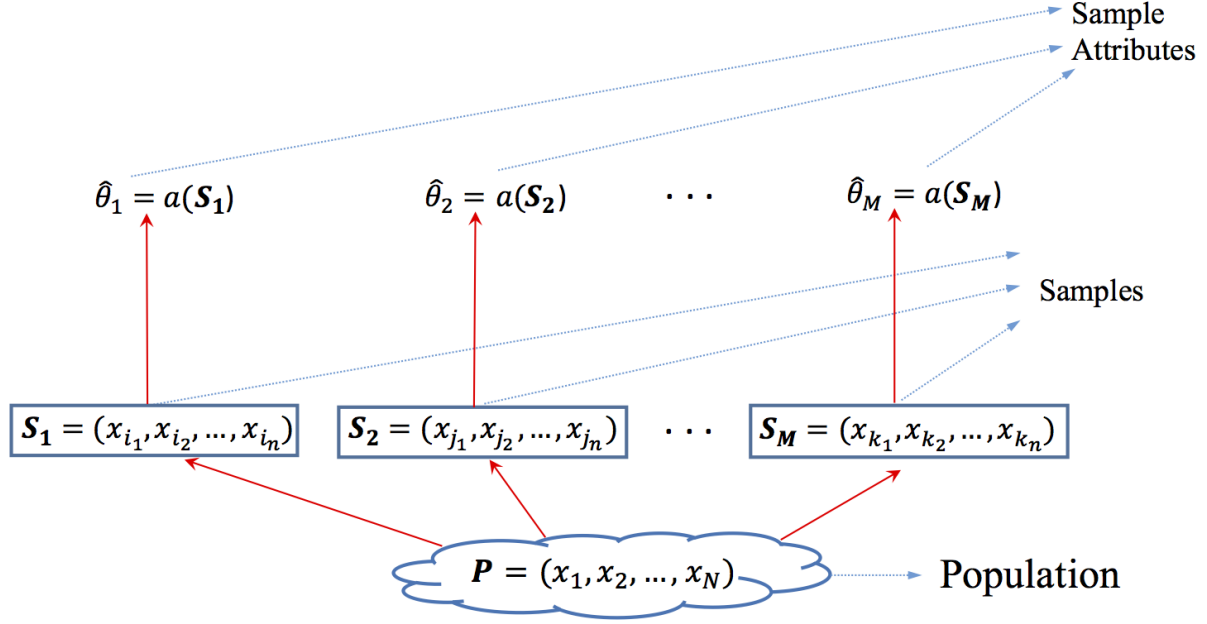


Figure 1:

- However, in practice we have only one sample.

Mimic Resampling

- We propose to mimic this process
 - by drawing B samples $\mathcal{S}_1^*, \dots, \mathcal{S}_B^*$ of size n independently from a population \mathcal{P}^* .
- Ideally, \mathcal{P}^* will be the study population \mathcal{P} ,
 - but our sample estimate of \mathcal{P} is the sample \mathcal{S} , so we take $\mathcal{P}^* = \mathcal{S}$ or $\hat{\mathcal{P}} = \mathcal{S}$.
 - However, we could use any estimate of the study population.
- The sample population has only n units, so without replacement sampling mechanism will immediately exhaust the population. Therefore, we sample **with replacement**.
- Drawing B samples $\mathcal{S}_1^*, \dots, \mathcal{S}_B^*$ of size n from \mathcal{P}^* with replacement,
 - the sample attribute values $a(\mathcal{S}_1^*), \dots, a(\mathcal{S}_B^*)$ now provide the information on the sampling distribution of interest.

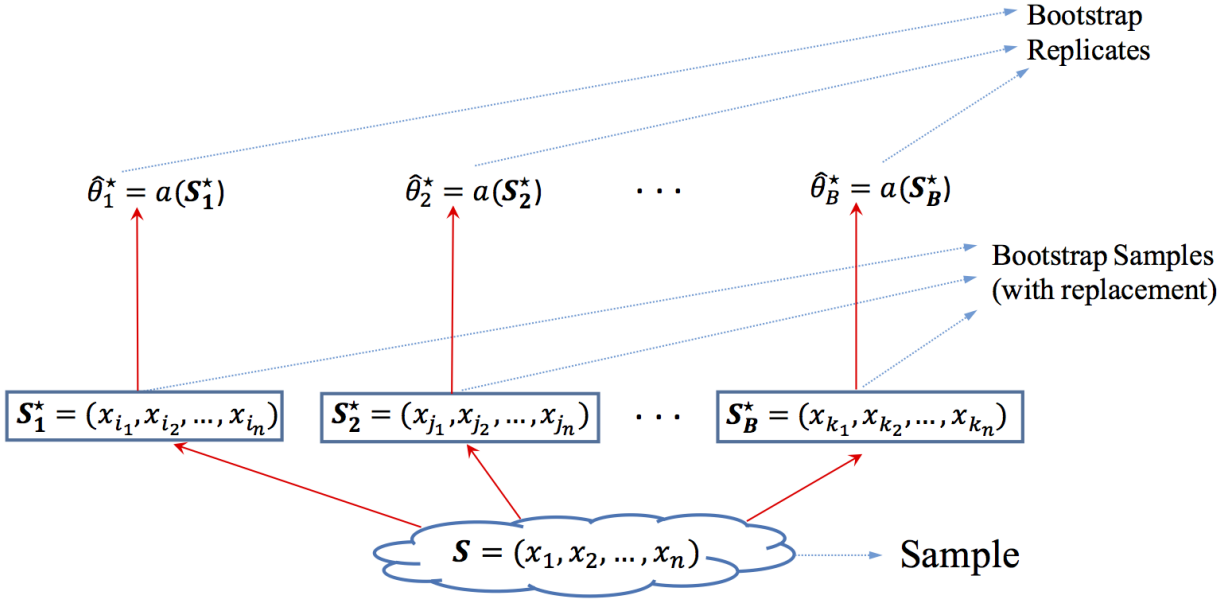


Figure 2:

The Bootstrap Method

- This approach of mimicking the sampling distribution was named the **bootstrap** method when it was first proposed in 1979 by Bradley Efron.
- The word “bootstrap” conveys the notion of starting something up from nothing as in “pulling oneself over a fence by one’s bootstraps”.
 - It suggests something for nothing, or something impossible to achieve.
- In Efron (1979), other possible names were *Swiss Army Knife*, *Meat Axe*, *Swan-Dive*, *Jack-Rabbit*.
- The Bootstrap path of Inductive inference




Figure 3:

Why is bootstrapping so important? What does it provide?

Shark Lengths Example

- Suppose we have a sample of $n = 6$ and we are interested in estimating the average shark length.
 - Do we expect the bootstrap to work if n is small?
 - The population has only $N = 65$ units so the sample size is a little under 10 percent of the population.
- We draw one sample \mathcal{S} drawn from \mathcal{P} using simple random sampling without replacement.

```
set.seed(341)
n <- 6
S <- getSample(1:65, n, replace = FALSE)
S

## [1] 10 58 24 4 50 46
```

- Then we draw B bootstrap samples from this single sample.
 - There are $n^n = 6^6 = 46656$ possible **bootstrap samples** of size $n = 6$ to select.
 - Here, we choose $B = 10000$ bootstrap samples $S_1^*, \dots, S_{10000}^*$:
 - We might choose a different B depending on the context.

```
Pstar <- S
B <- 10^4
set.seed(341)
Sstar <- sapply(1:B, FUN = function(b) getSample(Pstar, n, replace = T))
```

Here we then have a matrix bootstrap samples and the first bootstrap sample contains units

```
#dim(Sstar)
Sstar[,1]
```

```
## [1] 10 46 24 10 50 50
```

- We then compute whichever attribute might be of interest on each bootstrap sample.
 - e.g. we compute the length average for each bootstrap sample.

```
avesBootSamp <- sapply(1:B, FUN = function(i) mean(sharks[Sstar[,i], "Length"]))
#length(avesBootSamp)
```

The initial ten bootstrap sample averages are

```
round(avesBootSamp[1:10], 1)
```

```
## [1] 112.3 128.7 121.8 122.0 121.0 100.5 108.8 110.8 120.3 142.5
```

- The collection bootstrap averages is a population and
 - can be summarized like any other population
 - but is estimate of the sampling distribution of the sample average.

An estimate of the sampling distribution

- The distribution of any attribute over the bootstrap samples S_i^* from \mathcal{P}^* is a **bootstrap estimate** of the distribution of the same attribute over all possible samples S_i from \mathcal{P}

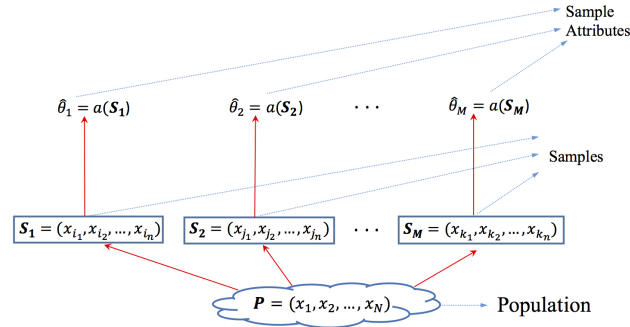


Figure 4:

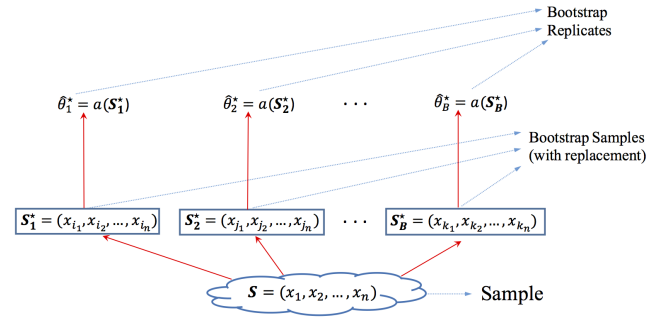


Figure 5:

Shark Length Comparision

- We can compare the **bootstrap estimate** to resampling from the population. To perform resampling from the population,




Figure 6: Bootstrap versus Sampling with $n=6$: averages

- First we have to draw $B = 10,000$ samples from the population and
- calculate the average on each sample.

```
avesSamp <- sapply(1:B, FUN = function(i) mean(sharks[sample(65, n), "Length"]))
#length(avesSamp)
```

The initial ten sample averages are

```
round(avesSamp[1:10], 1)
```

```
## [1] 184.3 160.5 157.2 166.2 178.2 167.0 165.3 138.8 144.3 131.3
```

```
savePar <- par(mfrow=c(1,2))
```

```
hPopAve <- hist(extendrange(c(avesSamp,avesBootSamp)), breaks = 50, plot = FALSE)
hist(avesBootSamp, xlim = range(avesSamp), breaks = hPopAve$breaks,
     freq = FALSE, col = "grey", main = "Bootstrap Averages", xlab="")
hist(avesSamp, xlim = range(avesSamp), breaks = hPopAve$breaks,
     freq = FALSE, col = "grey", main = "Sample Averages, n=6", xlab="")
```

- Note that the bootstrap estimate is, at best, as good as $a(\mathcal{S})$ (why?).
- As can be seen, the bootstrap distribution gives a sense of how an attribute varies.
 - e.g., to get some idea of the variability we can take the standard deviation of the bootstrap distribution.

- To easily compare the variability we can construct a histogram of the errors using

$$\begin{aligned}\text{sample error} &= a(\mathcal{S}) - a(\mathcal{P}) \\ \text{bootstrap sample error} &= a(\mathcal{S}^*) - a(\mathcal{S})\end{aligned}$$

```
savePar <- par(mfrow=c(1,2))
```

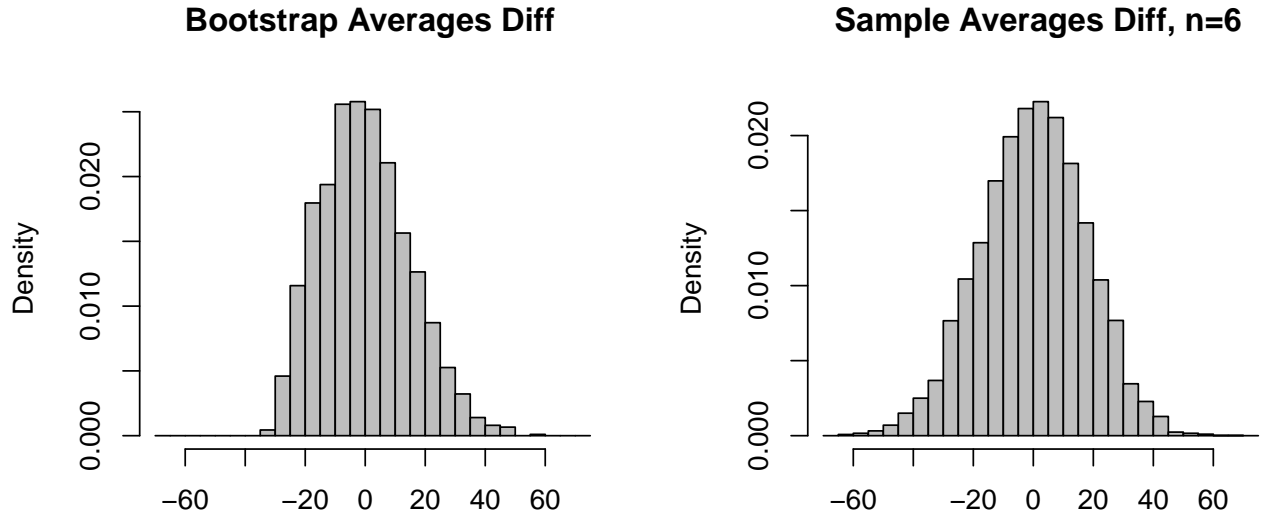


Figure 7: Bootstrap versus Sampling with n=6: averages

```
range.avediff <- extendrange(
  c(avesSamp- mean(avesSamp),
    avesBootSamp- mean(avesBootSamp)) )

hPopAvediff <- hist(range.avediff,
  breaks = 50, plot = FALSE)

hist(avesBootSamp - mean(avesBootSamp),
  xlim = range.avediff, breaks = hPopAvediff$breaks,
  freq = FALSE, col = "grey",
  main = "Bootstrap Averages Diff", xlab="")

hist(avesSamp - mean(avesSamp),
  xlim = range.avediff, breaks = hPopAvediff$breaks,
  freq = FALSE, col = "grey",
  main = "Sample Averages Diff, n=6", xlab="")
```

Bootstrap Standard Deviation

For any attribute, $a(\mathcal{P})$, the standard deviation of the corresponding sample attribute can be estimated from the bootstrap distributions with

$$\widehat{SD}_\star(\tilde{a}(\mathcal{S}^\star)) = \sqrt{\frac{\sum_{b=1}^B (a(\mathcal{S}_b^\star) - \bar{a}^\star)^2}{B-1}}$$

- where $\bar{a}^\star = \sum_{b=1}^B a(\mathcal{S}_b^\star)/B$ is the average of the attribute on the bootstrap samples $\mathcal{S}_1^\star, \dots, \mathcal{S}_B^\star$. Since B is usually large, it does not make any practical difference to use B or $B-1$ in the denominator of the standard deviation.

- This is an estimate of the standard deviation of the sampling distribution for the attribute $a(\mathcal{S})$.

The bootstrap estimate is

```
sdn(avesBootSamp)
```

```
## [1] 14.82945
```

whereas the standard deviation of the sample averages is

```
sdn(avesSamp)
```

```
## [1] 17.61667
```

The Average

- In the case of the arithmetic average $a(\mathcal{S}) = \sum_{u \in \mathcal{S}} y_u / n$
- The bootstrap estimate of the its standard deviation is

$$\widehat{SD}_*(\bar{Y}) = \sqrt{\frac{\sum_{b=1}^B (\bar{y}_b^* - \bar{y}^*)^2}{B-1}}$$

where $\bar{y}^* = \frac{\sum_{b=1}^B \bar{y}_b^*}{B}$.

- The standard deviation can also be estimated with

$$\widehat{SD}(\bar{Y}) = \frac{\hat{\sigma}}{\sqrt{n}} \sqrt{\frac{N-n}{N-1}} \quad \text{where} \quad \hat{\sigma} = \sqrt{\frac{\sum_{u \in \mathcal{S}} (y_u - \bar{y})^2}{n}}.$$

- Comparing these two estimates on the the length variable in shark encounters data, yields

```
n=6
N=dim(sharks)[1]
c(sdn(avesBootSamp), sdn(sharks[S, "Length"])/sqrt(n)*sqrt((N-n)/(N-1)))
```

```
## [1] 14.82945 14.31729
```

If the sample is not a good representative of the population, these two numbers might be quite different.

Comparing two estimators of the standard deviation

- We can compare the two estimators of standard deviation by a generating $m = 2000$ samples of size $n = 6$ and for each sample:
 - calculate $\frac{\hat{\sigma}}{\sqrt{n}} \sqrt{\frac{N-n}{N-1}}$ and then
 - generate $B = 200$ bootstrap samples then to obtain the bootstrap standard deviation.

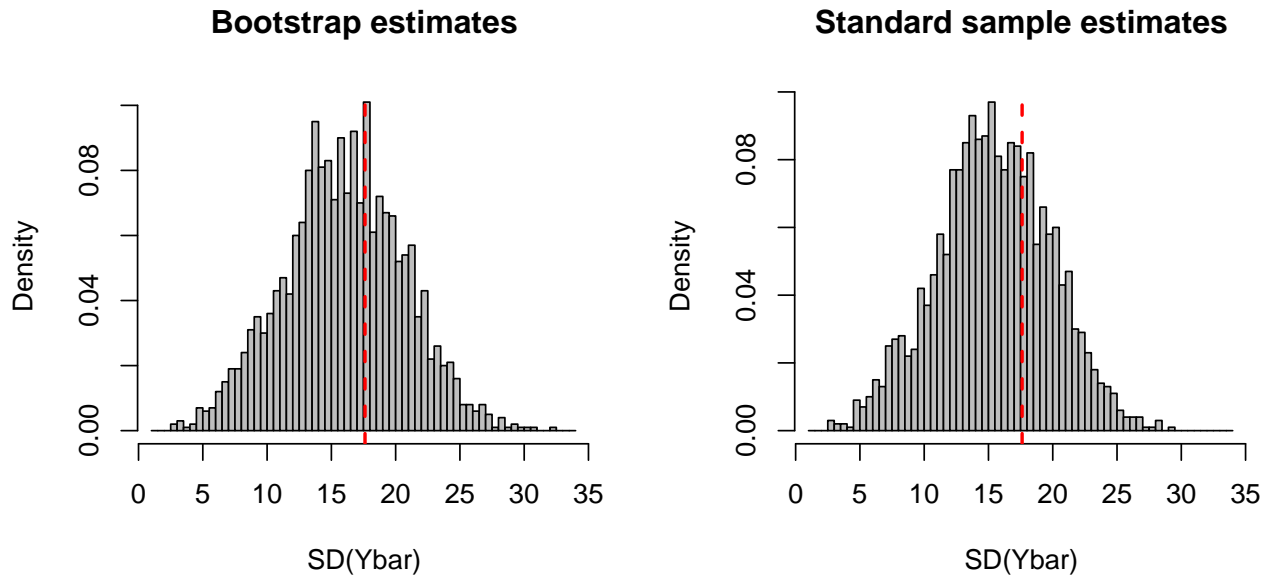


Figure 8: Comparing estimators of the standard deviation of the average

```
set.seed(341)
numSamps <- 2000
n <- 6

sampSes <- sapply(1:numSamps, FUN = function(i) getSample(popSharks, n))
B <- 200

valsBoot = t(apply(sampSes, 2, function(sam=NULL, B=NULL) {
  avesTemp = unlist(Map( function(i) {
    mean(sharks[sample(sam, n, replace = TRUE), "Length"]) }, 1:B) )

  c( sdn(avesTemp), sdn( sharks[sam, "Length"] )/sqrt(n)*sqrt((N-n)/(N-1)))
}, B=B ))

colnames(valsBoot) <- c("sdBoot", "sdStandard")

par(mfrow=c(1,2))
hTmp <- hist(extendrange(valsBoot), breaks = 50, plot = FALSE)
hist(valsBoot[, "sdBoot"], xlim = extendrange(valsBoot),
     breaks = hTmp$breaks, freq = FALSE, col = "grey",
     main = "Bootstrap estimates", xlab="SD(Ybar)")
abline(v=sdn(avesSamp), col="red", lty=2, lwd=2)

hist(valsBoot[, "sdStandard"], xlim = extendrange(valsBoot),
     breaks = hTmp$breaks, freq = FALSE, col = "grey",
     main = "Standard sample estimates", xlab="SD(Ybar)")
abline(v=sdn(avesSamp), col="red", lty=2, lwd=2)
```

- The two histograms are not too dissimilar in shape.
- We can further explore the differences between these two estimators.

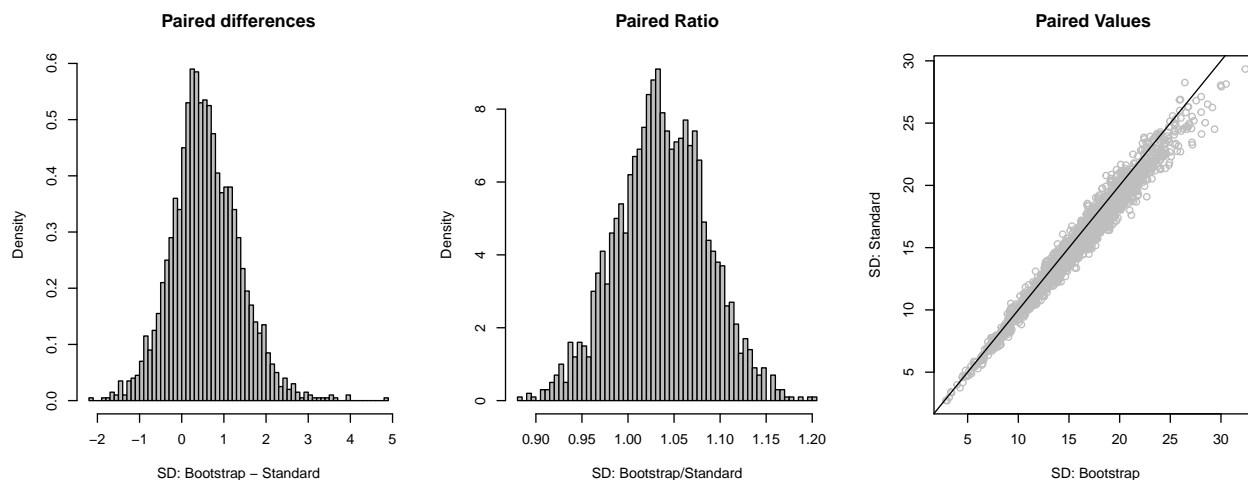


Figure 9: Paired differences in $SD(\bar{Y})$ estimates. The bootstrap estimator of standard deviation has produced, on average, slightly larger values than the standard.

- We can compare these two populations using the usual methods
 - Note: we have paired data because we calculated the two estimates on each sample.

```
par(mfrow=c(1,3))
hist(valsBoot[, "sdBoot"] - valsBoot[, "sdStandard"], main = "Paired differences", breaks=50, col="grey",
hist(valsBoot[, "sdBoot"] / valsBoot[, "sdStandard"], main = "Paired Ratio", breaks=50, col="grey", freq=
plot(valsBoot[, "sdBoot"], valsBoot[, "sdStandard"], main = "Paired Values ", col="grey", xlab="SD: Bootstrap",
abline(a=0, b=1)
```

- The third plot confirms that the bootstrap method may slightly over-estimate, the population SD.

Aside: n Versus $n - 1$

- In the above interval calculations the function `sdn(...)` was introduced and used which is an implementation of

$$\hat{\sigma} = \sqrt{\frac{\sum_{u \in \mathcal{S}} (y_u - \bar{y})^2}{n}}$$

which has n as a divisor. The built-in function `sd(...)` in R is an implementation of

$$\hat{\sigma} = \sqrt{\frac{\sum_{u \in \mathcal{S}} (y_u - \bar{y})^2}{n - 1}}$$

- For bootstrap interval calculations, a divisor of n is preferred.
 - This estimate has the advantage of being **replication invariant**.

- Replication invariant estimates are preferred and are often called **plug in estimates** in the bootstrap literature e.g. see Efron and Tibshirani (1994) .
- When n is reasonably large, there will be little practical difference between the two.

Estimating the sampling bias

- Recall that the **sampling bias**

$$\text{Sampling Bias} = E[a(\mathcal{S})] - a(\mathcal{P})$$

- We can use the bootstrap to estimate sampling bias via

$$\widehat{\text{sample bias}} = \text{average bootstrap sample error} = \bar{a}^* - a(\mathcal{S})$$

- where $\bar{a}^* = \sum_{b=1}^B a(\mathcal{S}_b^*)/B$ is the average of the attribute on the bootstrap samples $\mathcal{S}_1^*, \dots, \mathcal{S}_B^*$.

Biased Corrected Estimators

- Some estimators are unbiased and some are biased.
 - We would like to “correct” biased estimators.
 - i.e. make a biased estimator unbiased.
- If $a(\mathcal{S})$ was biased and we knew the bias then
 - we could add the correction to our attribute or estimator to make a new attribute $a^*(\mathcal{S})$ that is unbiased.

$$a^*(\mathcal{S}) = a(\mathcal{S}) - \text{Sampling Bias}$$

because

$$E[a^*(\mathcal{S})] = E[a(\mathcal{S}) - \text{Sampling Bias}] = E[a(\mathcal{S}) - E[a(\mathcal{S})] + a(\mathcal{P})] = a(\mathcal{P})$$

- Then the Bootstrap biased correct estimate is

$$a(\mathcal{S}) - \widehat{\text{bias}}[a(\mathcal{S})] = a(\mathcal{S}) - [\bar{a}^* - a(\mathcal{S})] = 2a(\mathcal{S}) - \bar{a}^*$$

Example with different Attributes

- Using the sample, 10, 58, 24, 4, 50, 46, we can estimate their bias and standard deviation of the corresponding estimators for the shark lengths:
 - median
 - standard deviation
 - skewness

The required functions and the samples estimates are

```
skew <- function(z) { 3*(mean(z) - median(z))/sdn(z) }

sam.len = sharks[S, "Length"]
sam.est = c(median(sam.len), sdn(sam.len), skew(sam.len))
round(sam.est,2)

## [1] 105.00 36.53 1.25
```

- Generating the bootstrap samples and calculating the attributes on each bootstrap sample

```
bootAttr <- apply(Sstar, 2, function(sam, len) {
  pop = len[sam]
  c( median(pop), sdn(pop), skew(pop) )
},
len = sharks[, "Length" ] )
# dim(bootAttr)
# the rows of boot Attr are the attributes (median,std. dev., and skewness)
## from 10,000 bootstrap samples.
```

The bootstrap estimate of the sampling distribution's standard deviation is

```
round(apply(bootAttr,1,sd),3)

## [1] 18.370 11.643 0.825
```

The bootstrap estimates of their sampling bias is

```
bias.est = apply(bootAttr,1,mean) - sam.est
round(bias.est,2)

## [1] 6.5 -5.3 -0.5
```

The biased corrected bootstrap estimates are:

```
round( sam.est - bias.est, 2)


## [1] 98.50 41.82 1.74

###
###
### or equivalently,
### round( 2*sam.est - apply(bootAttr,1,mean) ,2)
```

Estimates of their sampling distributions (no correction to the bias applied):

```
par(mfrow=c(1,3))
namHist = c("median", "standard deviation", "skewness")

for (i in 1:3) hist(bootAttr[i, ], main = namHist[i], xlab="Bootstrap Replicates",
  breaks=50, col="grey", freq=FALSE)
```



4.5.2 Bootstrap confidence intervals

- The bootstrap distribution provides a proxy for the sampling distribution for any sample attribute $a(\mathcal{S})$.
 - So we can use it to construct (at least approximate) confidence intervals for the unknown population attribute $a(\mathcal{P})$.

- If the bootstrap distribution is approximately normal then
 - noting that confidence intervals for sample averages, for example, have the following structure

$$\left[\bar{y} - c\widehat{SD}(\bar{Y}), \bar{y} + c\widehat{SD}(\bar{Y}) \right]$$

- Picking c such that $P(Z \leq c) = 1 - \alpha/2$ generates a $100(1 - \alpha)\%$ confidence interval (under Gaussian assumption).
- Rather than using the $\hat{\sigma}/\sqrt{n} \times \sqrt{(N - n)/(N - 1)}$ for $\widehat{SD}(\bar{Y})$,
 - we might use standard deviation of the bootstrap distribution of \bar{Y} to produce the standard deviation estimate $\widehat{SD}(\bar{Y})$.
 - The attraction of this approach, if it works, is that the same approach could be used for any attribute $a(\mathcal{S})$.

- A 95% **standard bootstrap interval** is

$$a(\mathcal{S}) \pm 1.96 \widehat{SD}_{boot}(a(\mathcal{S}))$$


where \widehat{SD}_{boot} is the bootstrap estimate of the standard deviation. What does 95% mean here?

- If we use the standard bootstrap interval then we should investigate its properties such as determining if we obtain the proper coverage probability.

Estimating the standard bootstrap interval coverage probability

- The average shark length of great whites was 151.86 inches (65 encounters).
- Suppose we are interested in a 95% confidence interval for the population mean.
- To get an estimate of the coverage probability for the bootstrap confidence interval, we generate 1000 samples of size $n = 6$ and for each sample
 - generate B bootstrap samples to obtain the bootstrap estimate of the standard deviation, $\widehat{SD}(\bar{Y})$, and
 - then construct the interval with $c = 1.96$.

100 individual 95% bootstrap confidence intervals



- As can be seen, only 86 of the 100 bootstrap intervals actually cover the average in the population.
 - This is a much lower coverage proportion than the expected 95 % suggested by using $c = 1.96$.

- Under-coverage is perhaps not unexpected as the c value is based on a Gaussian model which may not apply.

4.5.2.2 Percentile method

- The bootstrap distribution provides a proxy for the sampling distribution for any sample attribute $a(\mathcal{S})$.
 - So we can use it construct (at least approximate) confidence intervals for the unknown population attribute $a(\mathcal{P})$.
- If the bootstrap distribution is approximately a Gaussian distribution then confidence intervals inspired by the familiar structure

$$a(\mathcal{S}) \pm c \times \widehat{SD}(a(\mathcal{S}))$$
 - However, what if the bootstrap distribution is not approximately Gaussian?
- Then why not simply use the bootstrap distribution directly to construct a confidence interval?

Bootstrap versus Sampling Distribution

- Recall we compared
 - the Bootstrap distribution (based on sampling with replacement with $n = 6$ from a single sample) versus
 - Sampling Distribution (based on sampling without replacement with $n = 6$ from the population)
- Using the quantiles from sampling distribution we can obtain an interval that covers 95% of the sample averages.
 - This is (116.33, 184.83)
- Using the quantiles from bootstrap distribution we can obtain an estimate of the interval that covers 95% of the bootstrap averages.
 - This is (94.83, 151.83)
- The population average is 151.86

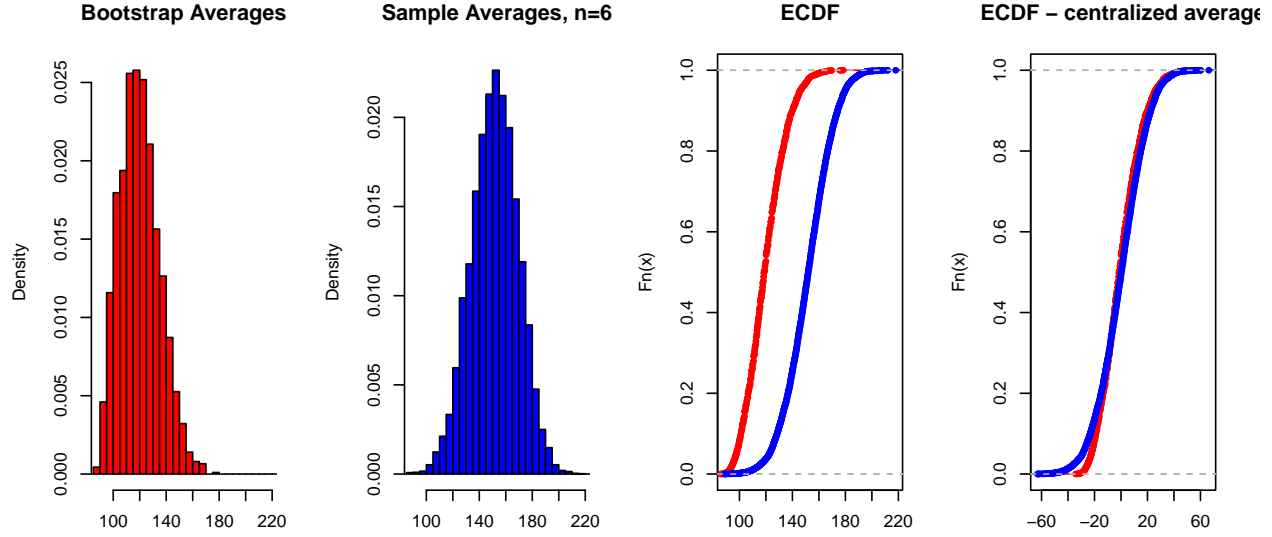


Figure 10: Bootstrap versus Sampling with $n=6$: averages

Percentile Method

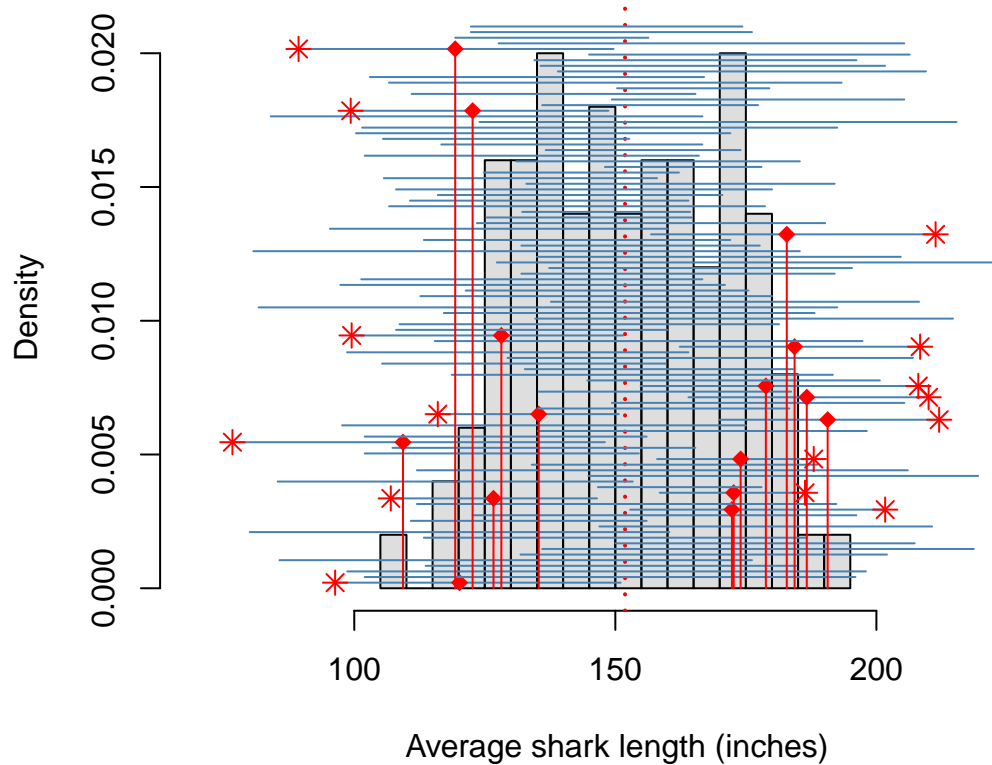
- For a given sample \mathcal{S} generate B bootstrap samples $\mathcal{S}_1^*, \dots, \mathcal{S}_B^*$
 - by sampling with replacement from the sample \mathcal{S} .
 - For the b^{th} bootstrap sample ($b = 1, \dots, B$), calculate $a_b = a(\mathcal{S}_b^*)$.
- From the values a_1, \dots, a_B , find
 - $a_{lower} = Q_a(p/2)$ and $a_{upper} = Q_a(1 - p/2)$
 - Then a $100(1 - p)\%$ confidence interval using the percentile method is $[a_{lower}, a_{upper}]$.
- This approach is **equivariant** to any one to one transformation (increasing function) of the attribute, say $T(a(\mathcal{P}))$.
 - i.e., the corresponding interval for $T(a(\mathcal{P}))$ is simply $[T(a_{lower}), T(a_{upper})]$!
 - So, we only need to determine the values a_{lower} and a_{upper} once for $a(\mathcal{P})$ and we have them for any $T(a(\mathcal{P}))$.
 - If the one to one transformation $T(a(\mathcal{P}))$ is monotonically decreasing, then the interval is $[T(a_{upper}), T(a_{lower})]$.

Assessing the coverage probability

- To estimate the percentile bootstrap interval coverage probability using a 95% confidence level.
- We generate 1000 samples of size $n = 6$ and for each sample $a_b = a(\mathcal{S}_b^*)$.
 - From the values a_1, \dots, a_B , find

- $a_{lower} = Q_a(p/2)$ and $a_{upper} = Q_a(1 - p/2)$
- Report the 95% percentile confidence intervals as $[a_{lower}, a_{upper}]$.

100 individual 95% bootstrap confidence intervals



- As can be seen, 85 of the 100 bootstrap intervals actually cover the average in the population.
 - This is a much lower coverage proportion than the expected 95 % suggested by using $c = 1.96$.

Comments

- Simplicity is the attraction of this method, and explains its continued popularity.
- This method is transformation invariant
- The coverage probability is often incorrect if the distribution of estimator is not nearly symmetric (Efron and Tibshirani, p. 178).
 - To improve the coverage probability we have to consider pivotal quantities or approximate quantities.

4.5.2.1 Bootstrap- t confidence intervals

- Some sampling experiments showed that

$$Z = \frac{\tilde{a}(\mathcal{S}) - a(\mathcal{P})}{\widetilde{SD}(\tilde{a}(\mathcal{S}))}$$

was approximately pivotal for $a(\mathcal{S}) = \bar{y}$

- its histogram (over all possible samples) was well approximated by a t -density.
 - This suggests instead of using a c from a Gaussian, we use a t -distribution on $n - 1$ degrees of freedom.
- Better still instead of approximate distribution (a t -distribution on $n - 1$ degrees of freedom) we use its actual distribution to determine the c value.
 - If $a(\mathcal{P})$ is the median or a measure of skewness, we would not expect the t -distribution to be a good approximation.
 - We have seen that the bootstrap is a method to approximate
 - the average error,
 - the standard deviation,
 - sampling distribution.
 - Here we will use the bootstrap to estimate the sampling distribution of

$$Z = \frac{\tilde{a}(\mathcal{S}) - a(\mathcal{P})}{\widetilde{SD}(\tilde{a}(\mathcal{S}))}$$

and use it to construct confidence intervals.

Sampling Distribution Approximation

- The t -distribution can approximate the sampling distribution of certain sample attributes.
 - This requires $\tilde{a}(\mathcal{S})$ to be approximately Gaussian over all possible samples and
 - an estimate of the standard deviation.
- The bootstrap distribution can approximate the sampling distribution of a sample attribute.
 - This requires computation.
 - However the bootstrap automatically adjusts its shape (and hence quantiles, etc.) to the form of $\tilde{a}(\mathcal{S})$.

- When we sample (say, $n = 6$) from all possible samples, we use the t -distribution to approximate

$$Z = \frac{\tilde{a}(\mathcal{S}) - a(\mathcal{P})}{\widehat{SD}(\tilde{a}(\mathcal{S}))}$$

```
ZPop <- (avesSamp - mean(sharks[, "Length"]))/sdn(avesSamp)
```

- To use the bootstrap to approximate this quantity replace
 - the sample \mathcal{S} with the bootstrap sample \mathcal{S}^* and
 - the population \mathcal{P} with an estimate \mathcal{P}^* , (the sample \mathcal{S}) and

$$Z^* = \frac{\tilde{a}(\mathcal{S}^*) - a(\mathcal{P}^*)}{\widehat{SD}(\tilde{a}(\mathcal{S}^*))} = \frac{\tilde{a}(\mathcal{S}^*) - a(\mathcal{S})}{\widehat{SD}(\tilde{a}(\mathcal{S}^*))}$$

```
ZBoot <- (avesBootSamp - mean(sharks[Pstar, "Length"]))/sdn(avesBootSamp)
```

The bootstrap- t

- For a given sample \mathcal{S}
 - calculate $a(\mathcal{S})$ and $\widehat{SD}(a(\mathcal{S}))$.
- Generate B bootstrap samples $\mathcal{S}_1^*, \dots, \mathcal{S}_B^*$ from \mathcal{S} and for each bootstrap sample:
 - Calculate $a(\mathcal{S}_b^*)$, $\widehat{SD}(a(\mathcal{S}_b^*))$ and

$$z_b = \frac{a(\mathcal{S}_b^*) - a(\mathcal{S})}{\widehat{SD}(a(\mathcal{S}_b^*))}.$$


- From the values z_1, \dots, z_B , find
 - $c_{lower} = Q_z(p/2)$ and $c_{upper} = Q_z(1 - p/2)$
- Then a $100(1 - p)\%$ bootstrap- t confidence interval is

$$\left[a(\mathcal{S}) - c_{upper} \times \widehat{SD}(a(\mathcal{S})), a(\mathcal{S}) - c_{lower} \times \widehat{SD}(a(\mathcal{S})) \right].$$

– **Note** the signs and positions of the constants c_{lower} and c_{upper} in the interval definition.

- This method (so far) requires an analytic form to calculate $\widehat{SD}(a(\mathcal{S}_b^*))$, which may or may not be available.

Sampling Distribution Comparison




We use the bootstrap distribution of T to find values c_{lower} and c_{upper} such that

$$Pr(c_{lower} \leq Z \leq c_{upper}) = (1 - p)$$

with $(1 - p)$ being the intended coverage probability.

100 individual 95% bootstrap confidence intervals



- Now 97 of the 100 bootstrap intervals cover the average in the population.
 - These intervals will be wider due to the bootstrap estimates of c_{lower} and c_{upper} than earlier intervals.

A thought on Standard Deviations

- When $a(\mathcal{S}) = \bar{y}$,
 - we have an analytic form for its standard deviation, namely

$$SD(\bar{Y}) = \frac{\sigma}{\sqrt{n}} \times \sqrt{\frac{N-n}{N-1}}$$

- Replacing σ by $\hat{\sigma}$ gives an estimate $\widehat{SD}(\bar{Y})$ based on the sample values y_u for $u \in \mathcal{S}$.

- Often, an analytic solution is not available for $\widehat{SD}(a(\mathcal{S}))$
 - An estimate can be obtained by using the bootstrap,
 - The standard deviation of the bootstrap values $a(\mathcal{S}_1^*), \dots, a(\mathcal{S}_B^*)$.
- However, in the bootstrap- t , we need an estimate of $SD(a(\mathcal{S}_b^*))$ for each bootstrap sample \mathcal{S}_b^* .
 - As an analytic form may not be available for $a(\mathcal{S})$, how can we get an estimate?

A Double Double

- **The Double Bootstrap** is general solution which applies the bootstrap method to the bootstrap sample \mathcal{S}_b^* .
- To apply a **bootstrap within a bootstrap**, for each bootstrap sample \mathcal{S}_b^*
 - we calculate a bootstrap estimate of the standard deviation by generating D bootstrap samples, $\mathcal{S}_1^{**}, \dots, \mathcal{S}_D^{**}$, from \mathcal{S}_b^* .
- Just as with the original bootstrap, this additional bootstrap adds another link in the chain of the inductive path as shown below.




Figure 11:

Example

- The `bootstrap_t_interval` function (see below), requires an attribute function `a` which will calculate $a(S)$ for any $S \subset \mathcal{P}$ as an argument.

```
bootstrap_t_interval <- function(S, a,
                                confidence = 0.95,
                                B = 1000, D = 30){
  ## Here S is the sample, a is a scalar-valued function a(S) of a sample S
  ## which returns the value for S of that attribute of interest
  ## confidence is the level of confidence
  ## B is the outer bootstrap count of replicates used to
  ## calculate the lower and upper limits
  ## D the inner bootstrap count of replicates used to
  ## estimate the standard deviation of the sample attribute
  ## for each (outer) bootstrap sample
  Pstar <- S
  aPstar <- a(Pstar)
  ## get (outer) bootstrap values
  bVals <- sapply(
    1:B,
    FUN = function(b) {
      Sstar <- getSample(Pstar, sampleSize, replace = TRUE)
      aSstar <- a(Sstar)
      ## get (inner) bootstrap values to
      ## estimate the SD
      Pstarstar <- Sstar
      SD_aSstar <- sdn(
        sapply(1:D,
              FUN = function(d){
                Sstarstar <- getSample(Pstarstar, sampleSize, replace = TRUE)
                ## return the attribute value
                a(Sstarstar)
              })
      )
      z <- (aSstar - aPstar)/SD_aSstar
      ## Return the two values
      c(aSstar = aSstar, z = z)
    })
  SDhat <- sdn(bVals["aSstar",])
  zVals <- bVals["z",]
  ## Now use these zVals to get the lower and upper
  ## c values.
  cValues <- quantile(zVals,
                      probs = c((1 - confidence)/2, (confidence + 1)/2),
                      na.rm = TRUE)
  cLower <- min(cValues)
  cUpper <- max(cValues)
  interval <- c(lower = aPstar - cUpper * SDhat,
               middle = aPstar,
               upper = aPstar - cLower * SDhat)
  interval
}
```


- In the case of the average shark length, this can be written as

```
a <- function(S) {mean(sharks[S, "Length"])}
a
```

A single bootstrap- t 95% confidence interval is

```
S <- getSample(popSharks, size = 5, replace = TRUE)
```

```
bootstrap_t_interval(S, a)
```

```
##      lower      middle      upper
## 69.64893 139.40000 201.72852
```

Assessing the coverage probability

- As before, we can conduct an experiment to assess the coverage probability for these intervals
 - This experiment takes a fair time to complete

```
confidence <- 0.95
B <- 1000
set.seed(341)

intervals <- sapply(1:numIntervals,
  FUN = function (i) {
    bootstrap_t_interval(samps[,i], a = a,
      confidence = confidence, B=B, D=30) } )
```

- Exactly 96 of the 100 bootstrap intervals cover the population attribute.

Comments

- Using the bootstrap- t intervals perform better than the percentile and the standard bootstrap intervals (based on the Gaussian distribution).
- According to [Efron and Tibshirani (1994), pp. 161-162] the bootstrap- t intervals are best suited for attributes which measure “location parameters” like the average, the median, a particular quantile, et cetera.
- For attributes that do not measure location, it may be necessary to first transform the attribute to a scale that produces bootstrap- t that have good coverage probabilities,
 - Or consider using the percentile method.

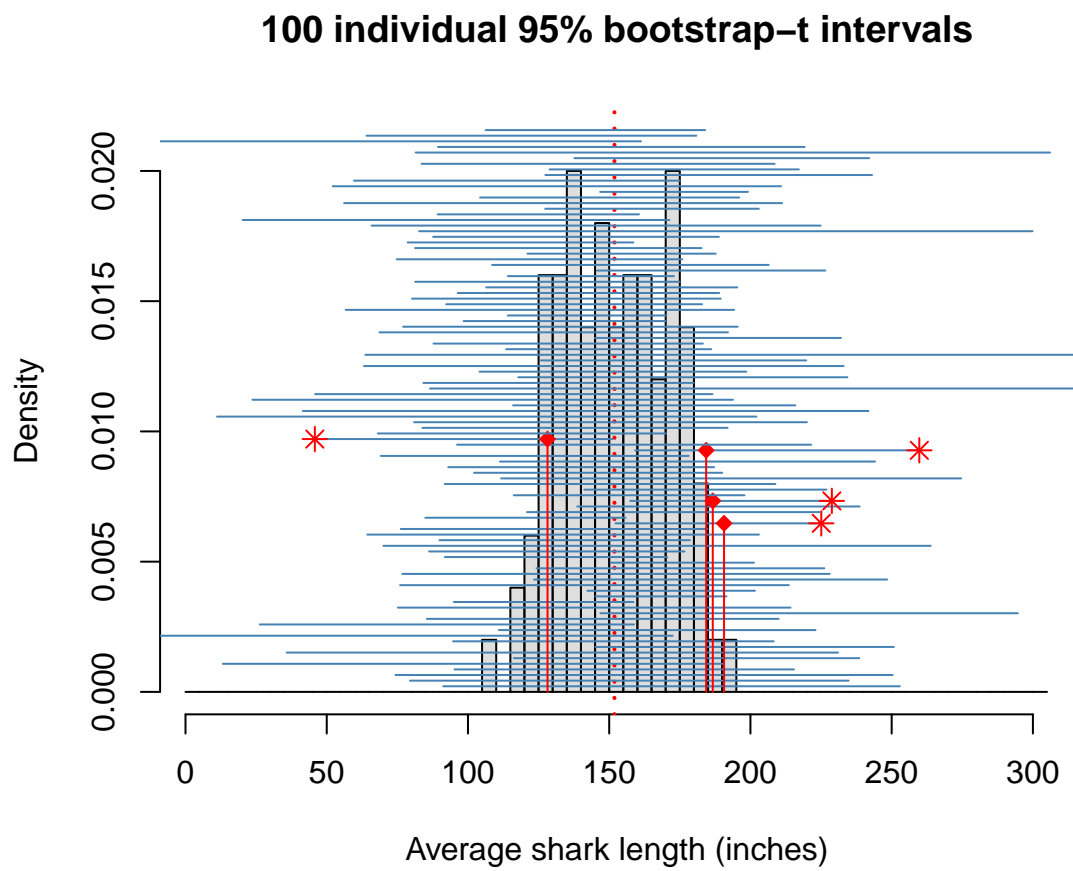


Figure 12: General bootstrap-t intervals

